PureSphere

The Thr'IF Musketeers

OT-7: Foundation of data engineering
Projet 2023 - Presentation


Tom DELAPORTE


Kevin KANAAN


Jorick PEPIN

Presentation plan

1. **PureSphere, what is it?**
   - 1.1. PureSphere reason to be
   - 1.2. Data sources
   - 1.3. Questions
2. **PureSphere overview**
   - 2.1. Conceptual view
   - 2.2. Logical view
3. **Ingestion**
4. **Staging**
   - 4.1. Cleansing
   - 4.2. Transformation
   - 4.3. Enrichment
5. **Production**
   - 5.1. Graph database
   - 5.2. Star Schema
6. **Further improvements**

1

PureSphere

What is it?

PureSphere reason to be

Have you ever wondered if we can we believe or not the statistics on industrial pollutant emissions?

→ **This is the starting point of PureSphere**

PureSphere is a **data pipeline** designed **to provide sufficient data** to **build trustworthy analysis** regarding the **impact of the industrial sites on their surrounding environments.**

Data sources

GÉORISQUES
Mieux connaître les risques sur le territoire

List of industrial facilities releasing pollutants.

Géorisques

• • •

GEOD'AIR

Reference data and statistics on air quality in France.

Geod'air

• • •

data.gouv.fr

hub'eau

List of physico chemical analysis of water quality.

Hub'eau

• • •

API

Questions

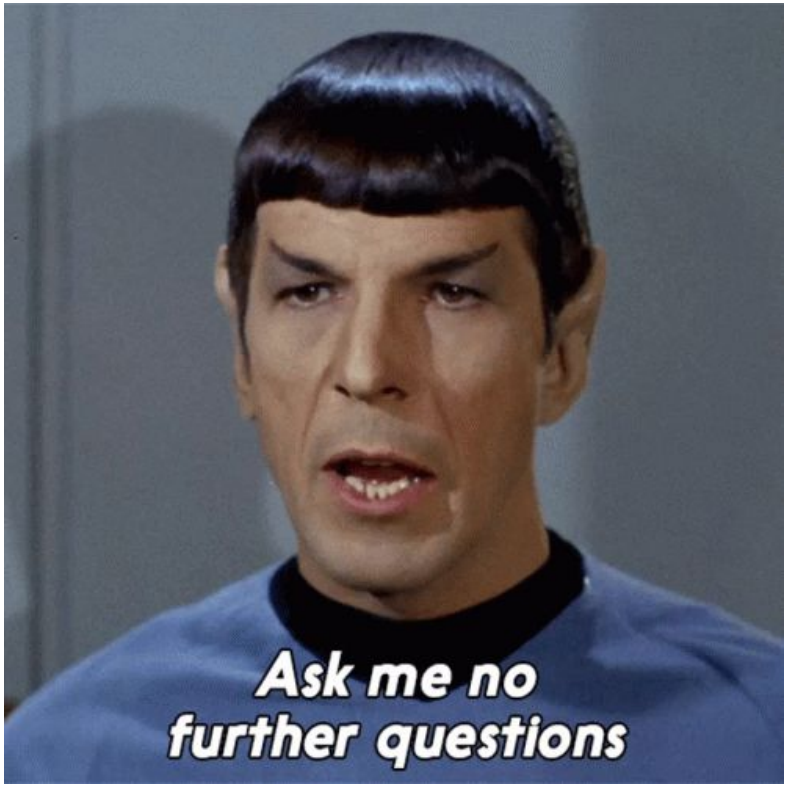

Easy

What are the zones for which we have information about the air quality and the water quality?

Medium

Can we see the impact of industrial sites on their surrounding area in terms of air and water quality?

Hard

Do the data from the air and water quality sensors coincide with the pollutant discharges given by Géorisques?

2

PureSphere

Overview

Conceptual view

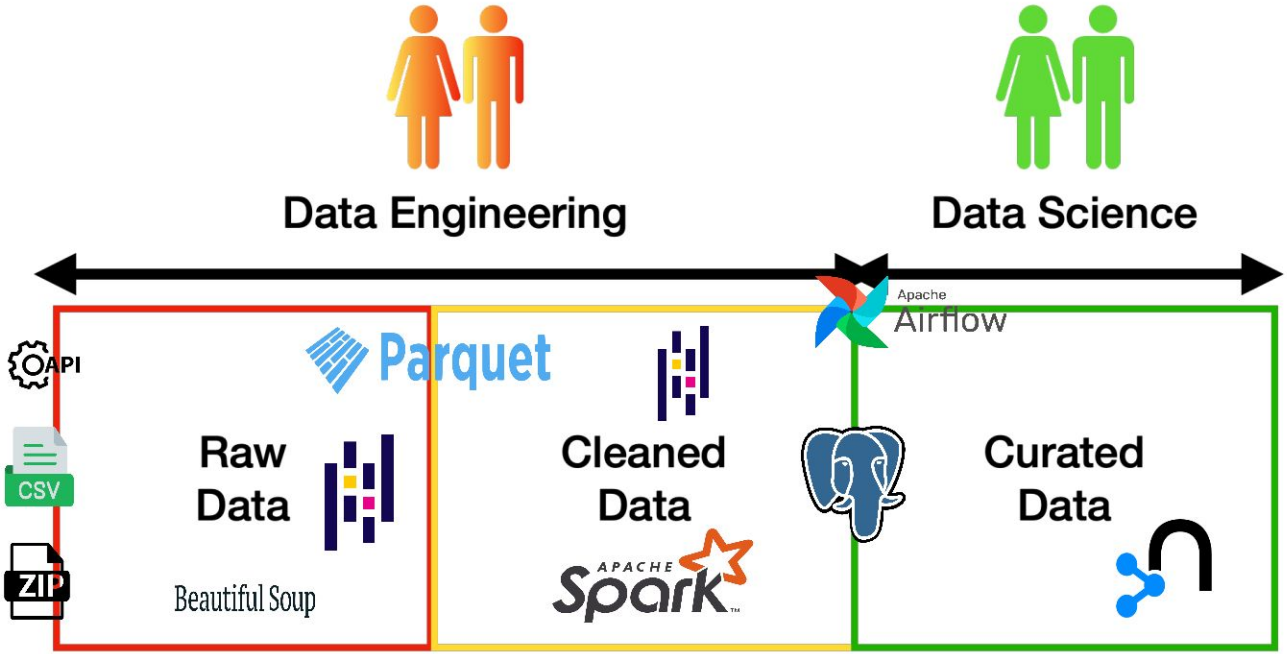

Logical view

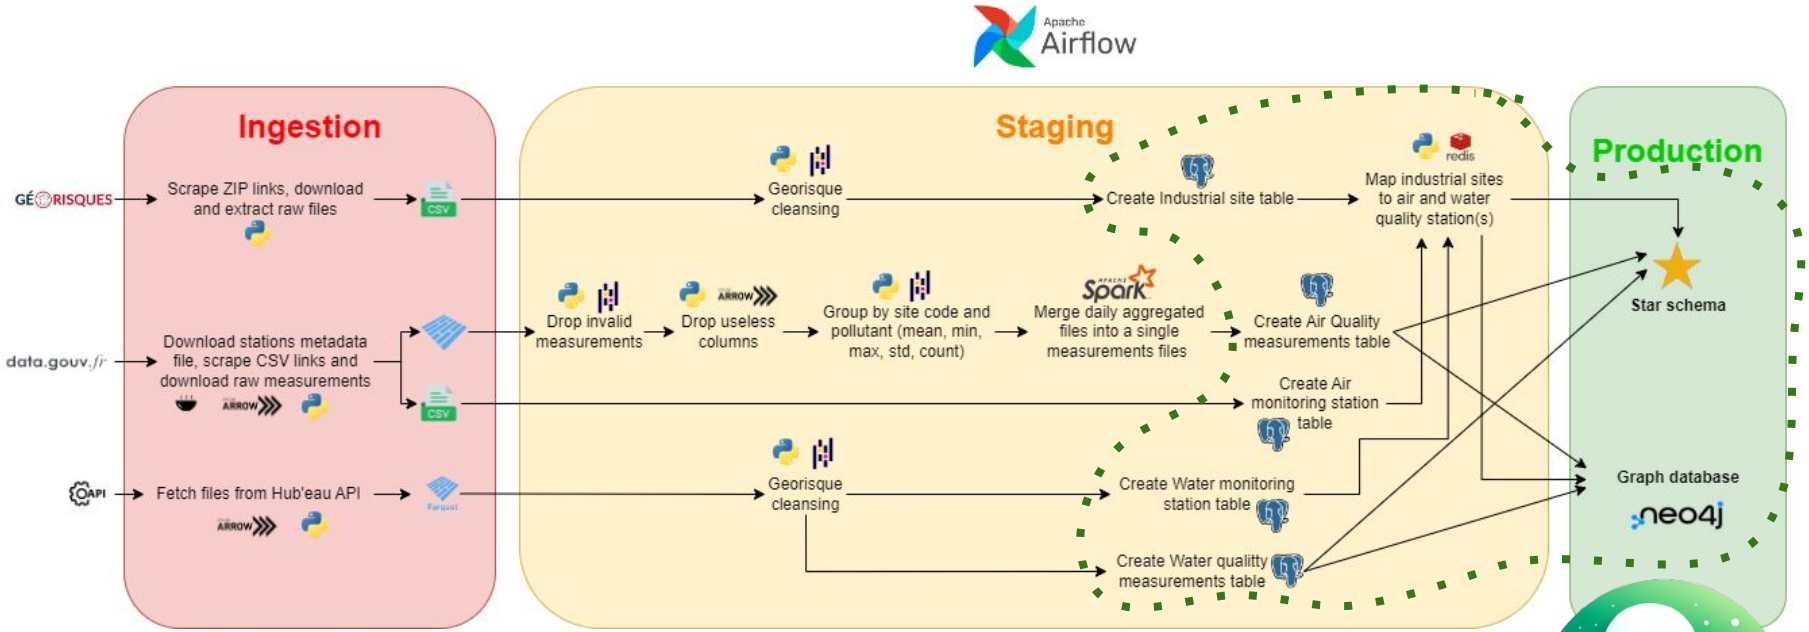

3. Ingestion

Several issues

- **APIs limitations**: ✖ OpenAQ → Limited by the call rate ~ Switch to Geod'air
  ✖ Hub'eau → Limited to 20000 rows/response ~ Smart requests

- **Data volume**: Geod'air example

≃3.5 GB / year 

≃66 MB / year

- **Sources aren't directly accessible**: ~ Extract link from web pages (Beautiful Soup)

4

PureSphere

Wrangling

Cleansing

I - Georisques

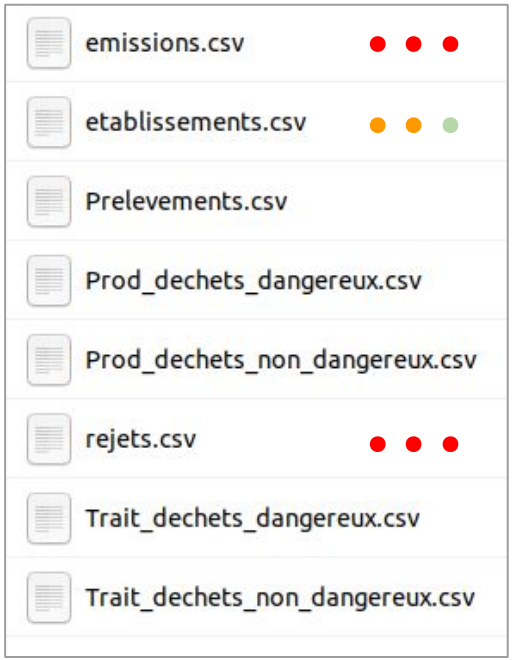

- Remove unnecessary data (files and columns)
- Transform literal addresses into GPS coordinates (latitude, longitude)

only 0.8% of valid values

GeoPy

coordonnees_x	coordonnees_y
180912.000000	6844188.000000
5.899579	45.588196
1175054.000000	6109594.000000
1176329.000000	6116976.000000
308717.000000	564590.000000
633535.000000	6904557.000000
422583.000000	6593363.000000
2.876031	47.508287
506855.000000	6291383.000000

Cleansing

II - Geod'air

After running the ingestion pipeline, the landing zone contains a Parquet file per day from the 1st of January 2021 until today. We also have a file containing relevant metadata about air monitoring stations.

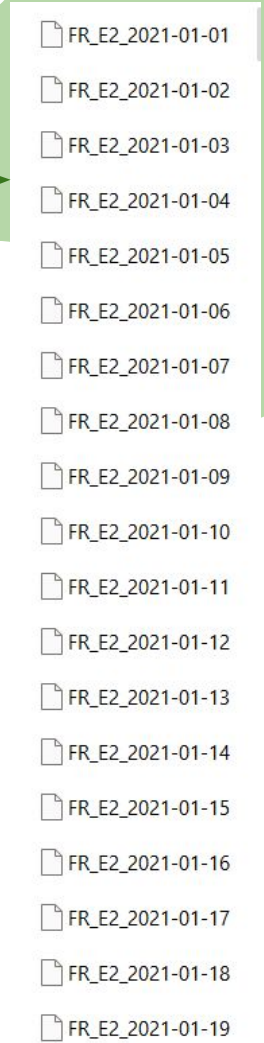

2 majors steps

pandas

1. **Remove invalid rows**: Some measurements are tagged as falsy or incoherent, some other have missing measurement. We kept data only coming from trustworthy stations.

2. **Drop useless columns**: For the sake of our project, we won't need all the columns contained by Geod'air files.

Cleansing

III - Hub'eau

In the Landing zone we also have a parquet file containing all the data that were measured during a year.

- Removed unnecessary columns
- Removed invalid rows

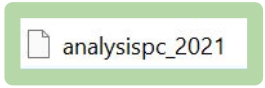

Transformation

Geod'air

A file per day containing hourly averaged pollutant concentrations for all the stations (a single station can monitor more than one pollutant).

For each file, we group by station and pollutant type and aggregate the following way:

- Average pollutant concentrations
- Min and max hourly average concentration recorded
- Standard deviation
- # measurements

Single file resulting of the concatenation of all the daily files augmented with a date column.

Batch processing

pandas

Transformat

Apache Spar

Single file resu d with a date column.

Spark 3.5.0 **Spark Master a**

URL: spark://172.21.0.4:7077
Alive Workers: 1
Cores in use: 3 Total, 0 Used
Memory in use: 4.0 GiB Total, 0.0 B Used
Resources in use:
Applications: 0 Running, 1 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers (1)

Worker Id				Resources
worker-20231121200404-172.21.0.5-34347				

Running Applications (0)

Application ID	Name			Duration

Completed Applications (1)

Application ID				Duration
app-20231121200530-0000			D	1.9 min


```
[2023-11-21, 20:05:57 UTC]                                    Manager: Finished task 1.0 in st
0.0 (TID 1) in 945 ms on 172.21.0.5 (executor 0) (1/1033)
```

Transformation

Hub'eau

Adding aggregations and calculating the average concentration of a pollutant for a day.

The goal is to have a file containing averaged pollutant concentrations for all the stations. In a way that matches for sure what we have with the air quality.

Batch processing

pandas

Transformation

Map industrial sites to their surrounding air and water quality monitoring stations based on:

- Spatial distance
- Kind of pollutant released by the industrial site VS kind of pollutant monitored by the station
- Any other relevant information we might have access to

The mapping represents heavy computations.

To avoid doing the computation twice, **we might use Redis to cache the mapping.**

Key	Value (List)
Industrial_site_1	[AQ_station_1, … WQ_station_42]

5

PureSphere

Production

Graph database

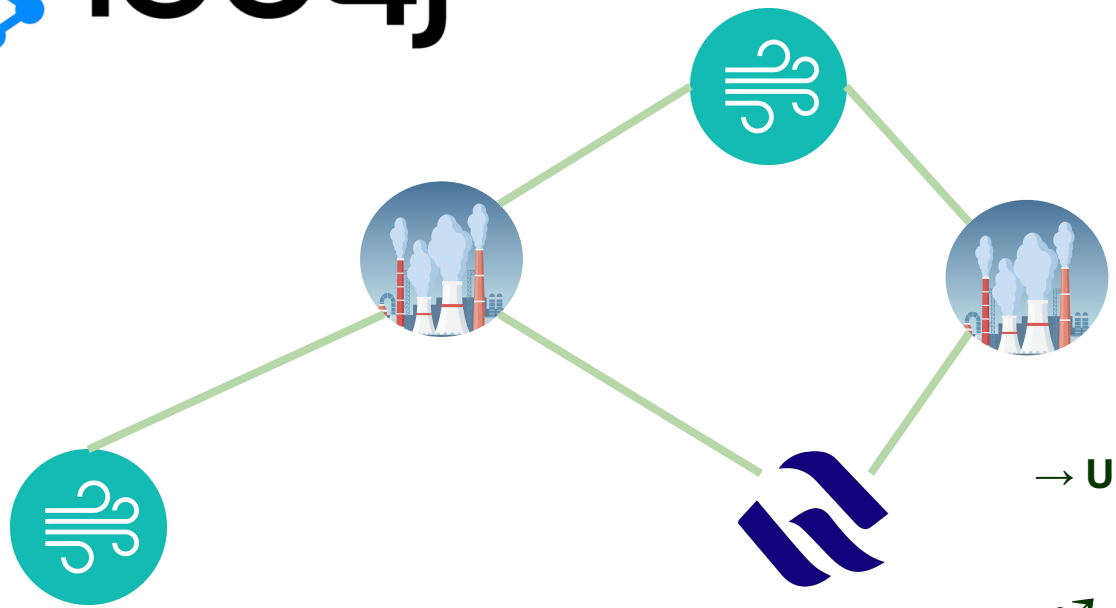

→ **Using Neo4J spatial functions**

↗ Hard • • •

Star schema

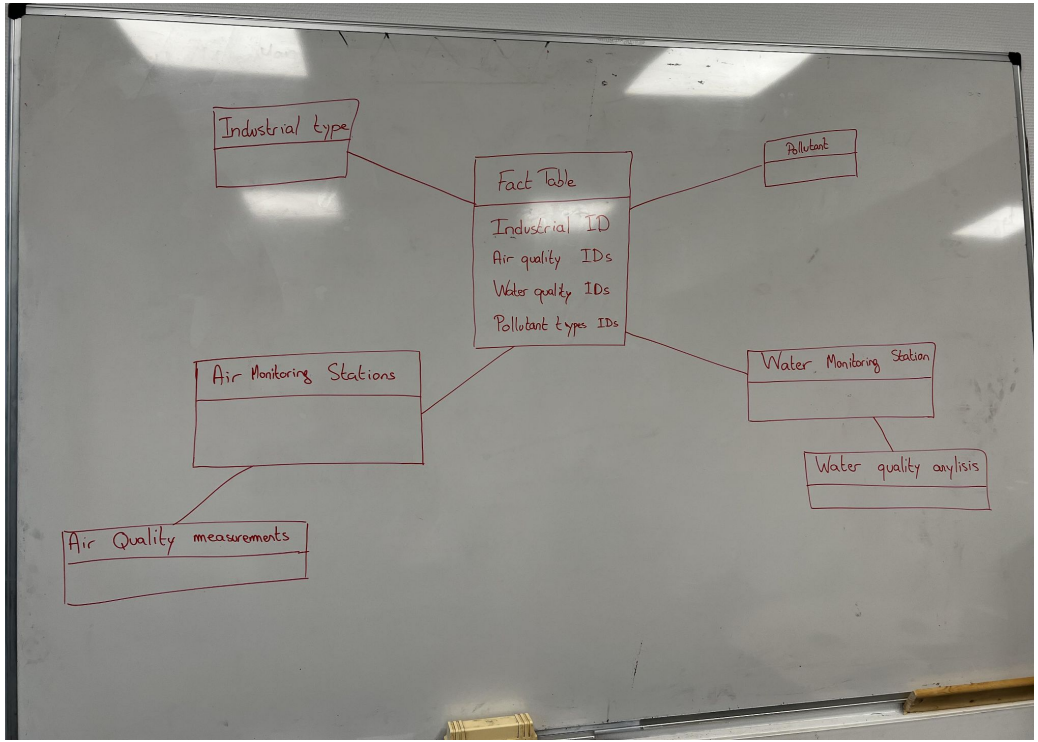

Easy

Medium

6

PureSphere

Further improvements

Further improvements

Use Redis for caching

Fine tune Spark to go even faster

Github Codespace

Thr'IF Muske

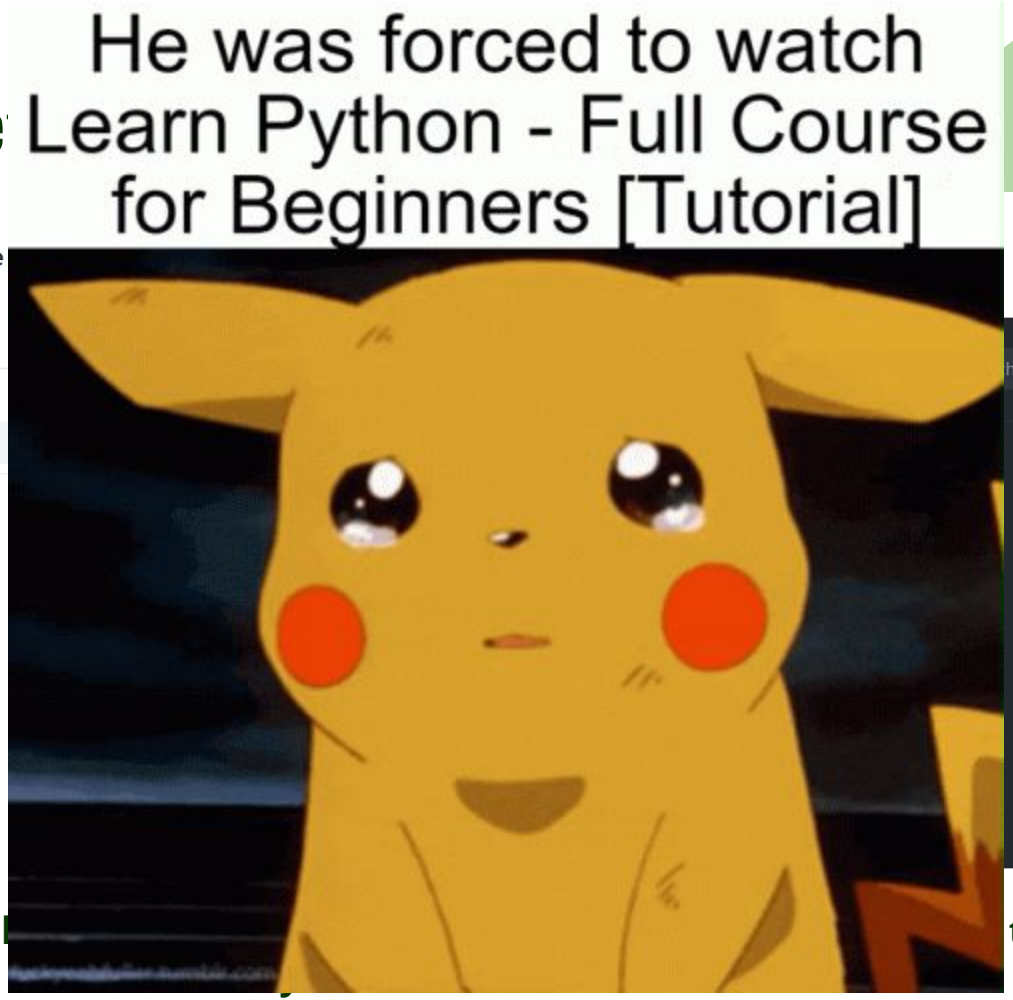

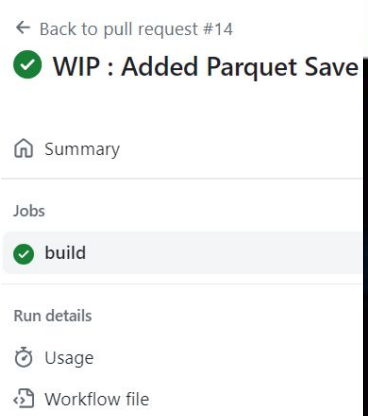

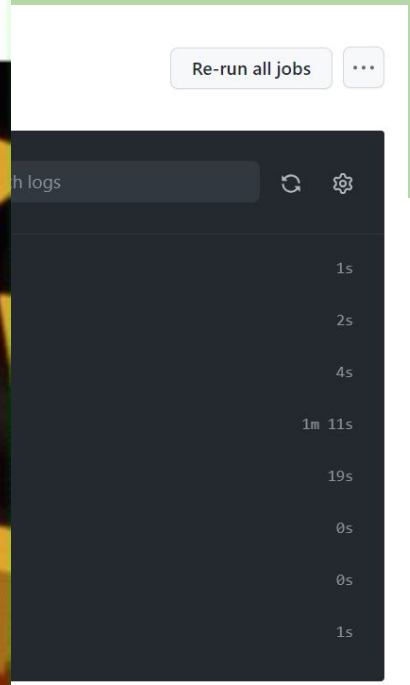

Automatic PyLint c **the PEP allergic**

CREDITS: This presentation template was created by **Slidesgo**, including icons by **Flaticon**, infographics & images by **Freepik** and illustrations by **Stories**

Thanks!

Do you have any questions?

The Thr'IF Musketeers

Tom DELAPORTE

Kevin KAANAN

Jorick PEPIN

PureSphere